Exact values and certified lower bounds for the no-three-in-line problem in the cube

Aleksei Kudriashov (Alex Komang)
Nusa Dua, Bali, Indonesia

August 2026 — version 0.1 (draft)

Abstract

Let $a(n)$ be the largest number of points of the grid $\{0, 1, \ldots, n-1\}^3$ no three of which are collinear. Pór and Wood proved $a(n) = \Theta(n^2)$; no exact values appear to have been recorded. We determine $a(n)$ for $n \leq 5$ by exhaustive search, certify lower bounds for $6 \leq n \leq 8$, and give an elementary construction showing $a(p) \geq p^2$ for every prime p. Every exact value was computed independently by two programs written to different designs, and every witness is re-verified by testing all triples in exact integer arithmetic.

1 Definitions and elementary bounds

Points of $G_n := \{0, \ldots, n-1\}^3$; a set is *lawful* if no three of its points are collinear in $\mathbb{R}^3$ (equivalently, all 2×2 minors of the difference matrix vanish for no triple). Each of the $3n^2$ axis-parallel lines carries at most two points of a lawful set, and each point lies on three of them, so $|S| \leq 2n^2$; this is the trivial upper bound, attained at $n = 2$.

Proposition 1. *For every prime p, $a(p) \geq p^2$.*

Proof. Let $Q(x, y) = x^2 - dy^2$ with d a quadratic non-residue modulo p, so that $Q(a, b) \equiv 0$ only for $a \equiv b \equiv 0$ (an anisotropic binary form; one exists for every p). Put $S = \{(x, y, Q(x, y) \bmod p) : 0 \leq x, y < p\}$, of size p^2. Suppose three of its points are collinear. Their projections to the (x, y)-plane are distinct (the third coordinate is a function of the first two), hence lie on a line with primitive integer direction (a, b), and the three points are $(x_0 + t_i a,\ y_0 + t_i b,\ \cdot\,)$ with distinct integers t_i, $|t_i| < p$. Along this line the third coordinate satisfies $z \equiv Q(x_0 + ta, y_0 + tb) = Q(a, b)\, t^2 + \ell(t) \pmod{p}$ with ℓ affine, whereas collinearity in $\mathbb{R}^3$ forces z to be an affine function of t. A quadratic congruence agreeing with an affine one at three points distinct modulo p has vanishing leading coefficient, so $Q(a, b) \equiv 0$ and therefore $(a, b) = (0, 0)$, a contradiction. $\square$

Neither $x^2 + y^2$ (anisotropic iff $p \equiv 3 \bmod 4$) nor $x^2 + xy + y^2$ (iff $p \equiv 2 \bmod 3$) works for all p; the choice $x^2 - dy^2$ does. The bound of Proposition 1 is not sharp: $a(5) \geq 40 > 25$ and $a(7) \geq 72 > 49$.

2 Method: a certified decision procedure

A line carrying at most two lattice points of $[n]^3$ forbids nothing, so

$$S \subseteq [n]^3 \text{ is admissible} \iff |S \cap L| \leq 2 \text{ for every line } L \text{ with } \geq 3 \text{ lattice points.} \tag{1}$$

This turns the question "is there an admissible set of size M?" into a Boolean satisfiability problem with n^3 variables — 125 for $n = 5$, 216 for $n = 6$ — one at-most-two cardinality constraint per rich line, and one global at-least-M constraint. The global counter is a totalizer truncated at $2n^2$; the truncation is legitimate because each of the n^2 axis-parallel lines carries at most two points, so $a(n) \le 2n^2$. We add the 47 lexicographic leader constraints of the cube group, which preserve satisfiability because the property in (1) is invariant under that group.

We adopted this route after measuring the alternative. A prefix-split exhaustive search on 42 cores closed 5 of 204 prefixes in two hours at $n = 5$, extrapolating to roughly 88 core-days; the same question is decided by a SAT solver in minutes. Speed, however, is not the reason. An exhaustive search establishes completeness by *its own* covering argument, which only its authors' code can check; an unsatisfiability answer comes with a DRAT certificate that an independent, widely used checker verifies without trusting our solver or our code at all.

Three separate links. It is worth being explicit about what is and is not certified, because the three claims below rest on different foundations.

1. *The formula expresses the problem.* Verified two independent ways: by comparing the set of triples forbidden by the rich-line constraints against the set with vanishing cross product, and by forcing each collinear triple as an assumption and requiring unsatisfiability. At $n = 4$ and $n = 5$ there are 376 and 1858 collinear triples; none was left unforbidden and none was forbidden spuriously. A DRAT checker cannot verify this link — it knows nothing of cubes and lines — and this is exactly where an encoding bug would hide.
2. *The formula is unsatisfiable.* Verified by `drat-trim` on the certificate emitted by `kissat`.
3. *A witness is admissible.* Verified by exhaustive integer arithmetic over all $\binom{|S|}{3}$ triples, by a program sharing no code with the search.

Links 1 and 2 give the upper bound; links 1 and 3 give the lower bound.

Calibration. Before computing anything new we reproduced every published value: $a(3) = 16$ and $a(4) = 28$, each with unsatisfiability at $M + 1$ in under a second, independently on two machines. For instances that resist a single solver call we split into independent pieces by fixing the contents of the first columns; the split is complete because three points of a column are collinear, so subsets of size at most two exhaust it. The aggregator refuses a verdict unless every piece of the manifest is present and unsatisfiable — a missing piece and an unsatisfiable one must never look alike.

3 Exact values

n	$a(n)$	$a(n)/n^2$	status
1	1	—	exhaustive
2	8	2.00	exhaustive (meets the trivial bound)
3	16	1.78	exhaustive
4	28	1.75	exhaustive
5	40	1.60	*lower bound certified; optimality in progress*

The exhaustive values were produced by two independent programs: a depth-first search over cells with per-line counters and column pruning, and an include/exclude search using, for every pair of cells, a precomputed mask of all cells collinear with them. Their node counts differ by a factor of about 72 ($17/3855/4.21 \cdot 10^7$ against $17/20068/3.03 \cdot 10^9$ for $n = 2, 3, 4$), so the agreement of the values is not an artefact of a shared design. For $n = 5$ the search is split into 256 independent prefixes (the subsets of size ≤ 2 taken in each of the first two z-columns; three points in a column would be collinear, so the prefixes cover the space without omission) and the completeness of the covering is checked mechanically before any claim of optimality.

4 Structure of the optima found

The witnesses produced for $n \leq 4$ are highly symmetric and share a pattern. Writing Stab for the stabiliser inside the group of the cube (order 48) and listing, for each axis, the number of points in the n layers perpendicular to it:

n	|Stab|	layer profile (identical for all three axes)	saturated axis lines
2	48	$4, 4$	12/12
3	24	$6, 4, 6$	24/27
4	8	$8, 6, 6, 8$	36/48

Each layer is a planar set with no three collinear, so it has at most $2n$ points; in every case the two outer layers attain that planar maximum and the inner ones have $2n-2$. The profile is the same along all three axes, and an axis line is called saturated when it carries two points.

We record this as an observation, not a law: fitting it gives $2n^2 - 2n + 4$, which matches $n \leq 4$ exactly and predicts $a(5) \geq 44$, whereas two independent searches (randomised destroy-and-repair, and CP-SAT) both plateau at 40, and a direct attempt to assemble a 44-point set from planar layers of profile $10, 8, 8, 8, 10$ failed in 4000 random trials. Either the pattern breaks at $n = 5$ or both searches are stuck; the exhaustive run in progress, which looks for any set of size ≥ 41, decides the question.

5 Certified lower bounds

n	$a(n) \geq$	$/n^2$
6	64	1.78
7	72	1.47
8	90	1.41

Obtained by randomised greedy search with destroy-and-repair; the value at $n = 7$ agrees with an independent CP-SAT run. Every configuration is re-verified by testing all $\binom{|S|}{3}$ triples with integer cross products. A dedicated exhaustive proof of optimality is out of reach for $n \geq 6$: a calibrated tree-size estimator (Knuth's random-probe method, validated against the true counts at $n = 4$ to within 4%) gives about $4.6 \cdot 10^{22}$ nodes for $n = 6$.

6 Reproducibility

All programs, journals, witnesses and the coverage checker are in the repository `https://github.com/iwasborninbali/saturation`.

AI/tool-use disclosure

The searches, the programs and the text were produced by autonomous AI agents (Anthropic Claude) under the direction of the author of record, who posed the question, chose what to compute, decided what to claim and takes responsibility for the content. Two agents worked with deliberately different program designs and checked each other's results; no value is reported here that was not obtained twice.